

chap01ex

September 15, 2021

1 Examples and Exercises from Think Stats, 2nd Edition

<http://thinkstats2.com>

Copyright 2016 Allen B. Downey

MIT License: <https://opensource.org/licenses/MIT>

```
[1]: from __future__ import print_function, division

import nsfg
```

1.1 Examples from Chapter 1

Read NSFG data into a Pandas DataFrame.

```
[2]: preg = nsfg.ReadFemPreg()
preg.head()
```

```
[2]:
```

| | caseid | pregordr | howpreg_n | howpreg_p | moscurrp | nowprgdk | pregend1 | \ |
|---|--------|----------|-----------|-----------|----------|----------|----------|---|
| 0 | 1 | 1 | NaN | NaN | NaN | NaN | 6.0 | |
| 1 | 1 | 2 | NaN | NaN | NaN | NaN | 6.0 | |
| 2 | 2 | 1 | NaN | NaN | NaN | NaN | 5.0 | |
| 3 | 2 | 2 | NaN | NaN | NaN | NaN | 6.0 | |
| 4 | 2 | 3 | NaN | NaN | NaN | NaN | 6.0 | |

| | pregend2 | nbrnaliv | multbrth | ... | laborfor_i | religion_i | metro_i | \ |
|---|----------|----------|----------|-----|------------|------------|---------|---|
| 0 | NaN | 1.0 | NaN | ... | 0 | 0 | 0 | |
| 1 | NaN | 1.0 | NaN | ... | 0 | 0 | 0 | |
| 2 | NaN | 3.0 | 5.0 | ... | 0 | 0 | 0 | |
| 3 | NaN | 1.0 | NaN | ... | 0 | 0 | 0 | |
| 4 | NaN | 1.0 | NaN | ... | 0 | 0 | 0 | |

| | basewgt | adj_mod_basewgt | finalwgt | secu_p | sest | cmintvw | \ |
|---|-------------|-----------------|--------------|--------|------|---------|---|
| 0 | 3410.389399 | 3869.349602 | 6448.271112 | 2 | 9 | NaN | |
| 1 | 3410.389399 | 3869.349602 | 6448.271112 | 2 | 9 | NaN | |
| 2 | 7226.301740 | 8567.549110 | 12999.542264 | 2 | 12 | NaN | |
| 3 | 7226.301740 | 8567.549110 | 12999.542264 | 2 | 12 | NaN | |
| 4 | 7226.301740 | 8567.549110 | 12999.542264 | 2 | 12 | NaN | |

```

    totalwgt_lb
0      8.8125
1      7.8750
2      9.1250
3      7.0000
4      6.1875

```

[5 rows x 244 columns]

Print the column names.

```
[3]: preg.columns
```

```
[3]: Index(['caseid', 'pregordr', 'howpreg_n', 'howpreg_p', 'moscurrp', 'nowprgdk',
          'pregend1', 'pregend2', 'nbrnaliv', 'multbrth',
          ...,
          'laborfor_i', 'religion_i', 'metro_i', 'basewgt', 'adj_mod_basewgt',
          'finalwgt', 'secu_p', 'sest', 'cmintvw', 'totalwgt_lb'],
          dtype='object', length=244)
```

Select a single column name.

```
[4]: preg.columns[1]
```

```
[4]: 'pregordr'
```

Select a column and check what type it is.

```
[5]: pregordr = preg['pregordr']
     type(pregordr)
```

```
[5]: pandas.core.series.Series
```

Print a column.

```
[7]: pregordr
```

```
[7]: 0      1
     1      2
     2      1
     3      2
     4      3
     5      1
     6      2
     7      3
     8      1
     9      2
    10      1
    11      1
    12      2
    13      3
    14      1
    15      2
    16      3
```

| | |
|-------|---|
| 17 | 1 |
| 18 | 2 |
| 19 | 1 |
| 20 | 2 |
| 21 | 1 |
| 22 | 2 |
| 23 | 1 |
| 24 | 2 |
| 25 | 3 |
| 26 | 1 |
| 27 | 1 |
| 28 | 2 |
| 29 | 3 |
| .. | |
| 13563 | 2 |
| 13564 | 3 |
| 13565 | 1 |
| 13566 | 1 |
| 13567 | 1 |
| 13568 | 2 |
| 13569 | 1 |
| 13570 | 2 |
| 13571 | 3 |
| 13572 | 4 |
| 13573 | 1 |
| 13574 | 2 |
| 13575 | 1 |
| 13576 | 1 |
| 13577 | 2 |
| 13578 | 1 |
| 13579 | 2 |
| 13580 | 1 |
| 13581 | 2 |
| 13582 | 3 |
| 13583 | 1 |
| 13584 | 2 |
| 13585 | 1 |
| 13586 | 2 |
| 13587 | 3 |
| 13588 | 1 |
| 13589 | 2 |
| 13590 | 3 |
| 13591 | 4 |
| 13592 | 5 |

Name: pregordr, Length: 13593, dtype: int64

Select a single element from a column.

```
[6]: pregordr[0]
```

```
[6]: 1
```

Select a slice from a column.

```
[7]: pregordr[2:5]
```

```
[7]: 2    1  
     3    2  
     4    3
```

Name: pregordr, dtype: int64

Select a column using dot notation.

```
[8]: pregordr = preg.pregordr
```

Count the number of times each value occurs.

```
[9]: preg.outcome.value_counts().sort_index()
```

```
[9]: 1    9148  
     2    1862  
     3     120  
     4    1921  
     5     190  
     6     352
```

Name: outcome, dtype: int64

Check the values of another variable.

```
[10]: preg.birthwgt_lb.value_counts().sort_index()
```

```
[10]: 0.0      8  
     1.0     40  
     2.0     53  
     3.0     98  
     4.0    229  
     5.0    697  
     6.0   2223  
     7.0   3049  
     8.0   1889  
     9.0    623  
    10.0    132  
    11.0     26  
    12.0     10  
    13.0      3  
    14.0      3  
    15.0      1
```

Name: birthwgt_lb, dtype: int64

Make a dictionary that maps from each respondent's caseid to a list of indices into the pregnancy DataFrame. Use it to select the pregnancy outcomes for a single respondent.

```
[11]: caseid = 10229  
     preg_map = nsfg.MakePregMap(preg)
```

```
indices = preg_map[caseid]
preg.outcome[indices].values
```

```
[11]: array([4, 4, 4, 4, 4, 4, 1], dtype=int64)
```

1.2 Exercises

Select the birthord column, print the value counts, and compare to results published in the [codebook](#)

```
[12]: preg.birthord.value_counts().sort_index()
```

```
[12]: 1.0    4413
      2.0    2874
      3.0    1234
      4.0     421
      5.0     126
      6.0      50
      7.0     20
      8.0       7
      9.0        2
     10.0        1
      Name: birthord, dtype: int64
```

We can also use `isnull` to count the number of nans.

```
[13]: preg.birthord.isnull().sum()
```

```
[13]: 4445
```

Select the prglngth column, print the value counts, and compare to results published in the [codebook](#)

```
[14]: preg.prglngth.value_counts().sort_index()
```

```
[14]: 0      15
      1       9
      2      78
      3     151
      4     412
      5     181
      6     543
      7     175
      8     409
      9     594
     10     137
     11     202
     12     170
     13     446
     14      29
     15      39
     16      44
     17     253
```

```

18      17
19      34
20      18
21      37
22     147
23      12
24      31
25      15
26     117
27        8
28      38
29      23
30     198
31      29
32     122
33      50
34      60
35     357
36     329
37     457
38     609
39    4744
40    1120
41     591
42     328
43     148
44      46
45      10
46        1
47        1
48        7
49        2
Name: prglngth, dtype: int64

```

To compute the mean of a column, you can invoke the mean method on a Series. For example, here is the mean birthweight in pounds:

```
[15]: preg.totalwgt_lb.mean()
```

```
[15]: 7.265628457623368
```

Create a new column named `totalwgt_kg` that contains birth weight in kilograms. Compute its mean. Remember that when you create a new column, you have to use dictionary syntax, not dot notation.

```
[16]: totalwgt_kg = preg['totalwgt_lb']/2.2046
```

`nsfg.py` also provides `ReadFemResp`, which reads the female respondents file and returns a `DataFrame`:

```
[17]: resp = nsfg.ReadFemResp()
```

`DataFrame` provides a method `head` that displays the first five rows:

```
[18]: resp.head()
```

```
[18]:   caseid  rscrinf  rdormres  rostscrn  rscreenhisp  rscreenrace  age_a  \
0    2298        1         5         5           1           5.0    27
1    5012        1         5         1           5           5.0    42
2   11586        1         5         1           5           5.0    43
3    6794        5         5         4           1           5.0    15
4     616        1         5         4           1           5.0    20

   age_r  cmbirth  agescrn  ...  pubassis_i  basewgt  adj_mod_basewgt  \
0     27     902      27  ...           0  3247.916977    5123.759559
1     42     718      42  ...           0  2335.279149    2846.799490
2     43     708      43  ...           0  2335.279149    2846.799490
3     15    1042      15  ...           0  3783.152221    5071.464231
4     20     991      20  ...           0  5341.329968    6437.335772

   finalwgt  secu_r  sest  cmintvw  cmlstyr  screentime  intvlngh
0  5556.717241      2   18    1234    1222    18:26:36  110.492667
1  4744.191350      2   18    1233    1221    16:30:59   64.294000
2  4744.191350      2   18    1234    1222    18:19:09   75.149167
3  5923.977368      2   18    1234    1222    15:54:43   28.642833
4  7229.128072      2   18    1233    1221    14:19:44   69.502667
```

```
[5 rows x 3087 columns]
```

Select the age_r column from resp and print the value counts. How old are the youngest and oldest respondents?

```
[23]: resp.age_r.value_counts().sort_index()
```

```
[23]: 15    217
16    223
17    234
18    235
19    241
20    258
21    267
22    287
23    282
24    269
25    267
26    260
27    255
28    252
29    262
30    292
31    278
32    273
33    257
```

```

34    255
35    262
36    266
37    271
38    256
39    215
40    256
41    250
42    215
43    253
44    235

```

```
Name: age_r, dtype: int64
```

We can use the caseid to match up rows from resp and preg. For example, we can select the row from resp for caseid 2298 like this:

```
[19]: resp[resp.caseid==2298]
```

```

[19]:   caseid  rscrinf  rdormres  rostscrn  rscreenhisp  rscreenrace  age_a \
0    2298         1         5         5             1          5.0    27

      age_r  cmbirth  agescrn  ...  pubassis_i      basewgt  adj_mod_basewgt \
0      27     902      27  ...           0  3247.916977    5123.759559

      finalwgt  secu_r  sest  cmintvw  cmlstyr  screentime  intvlngh
0  5556.717241         2   18    1234    1222    18:26:36   110.492667

```

```
[1 rows x 3087 columns]
```

And we can get the corresponding rows from preg like this:

```
[20]: preg[preg.caseid==2298]
```

```

[20]:   caseid  pregordr  howpreg_n  howpreg_p  moscurrp  nowprgdk  pregend1 \
2610    2298         1        NaN        NaN        NaN        NaN        6.0
2611    2298         2        NaN        NaN        NaN        NaN        6.0
2612    2298         3        NaN        NaN        NaN        NaN        6.0
2613    2298         4        NaN        NaN        NaN        NaN        6.0

      pregend2  nbrnaliv  multbrth  ...  laborfor_i  religion_i  metro_i \
2610        NaN        1.0        NaN  ...         0           0          0
2611        NaN        1.0        NaN  ...         0           0          0
2612        NaN        1.0        NaN  ...         0           0          0
2613        NaN        1.0        NaN  ...         0           0          0

      basewgt  adj_mod_basewgt      finalwgt  secu_p  sest  cmintvw \
2610  3247.916977    5123.759559  5556.717241         2   18      NaN
2611  3247.916977    5123.759559  5556.717241         2   18      NaN
2612  3247.916977    5123.759559  5556.717241         2   18      NaN
2613  3247.916977    5123.759559  5556.717241         2   18      NaN

```


| | totalwgt_lb |
|------|-------------|
| 2610 | 6.8750 |
| 2611 | 5.5000 |
| 2612 | 4.1875 |
| 2613 | 6.8750 |

[4 rows x 244 columns]

How old is the respondent with caseid 1?

```
[21]: resp[resp.caseid==1].age_r
```

```
[21]: 1069    44
      Name: age_r, dtype: int64
```

What are the pregnancy lengths for the respondent with caseid 2298?

```
[24]: preg[preg.caseid==2298].prglngth
```

```
[24]: 2610    40
      2611    36
      2612    30
      2613    40
      Name: prglngth, dtype: int64
```

What was the birthweight of the first baby born to the respondent with caseid 5012?

```
[25]: preg[preg.caseid==5012].birthwgt_lb
```

```
[25]: 5515    6.0
      Name: birthwgt_lb, dtype: float64
```

```
[ ]:
```